

Module 9: Recommended Exercises

Statistical Learning V2021

alexaoh

16 mars, 2021

Problem 1

Work through the lab in Section 9.6.1 of ISLR.

Problem 2 (Book Ex.2)

We have seen that in $p = 2$ dimensions, a linear decision boundary takes the form $\beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0$. We now investigate a non-linear decision boundary.

a) Sketch the curve

$$(1 + X_1)^2 + (2 - X_2)^2 = 4.$$

Done by hand. This is a circle with center $(-1, 2)$ and radius 2.

b) On your sketch, indicate the set of points for which

$$(1 + X_1)^2 + (2 - X_2)^2 > 4,$$

as well as the set of points for which

$$(1 + X_1)^2 + (2 - X_2)^2 \leq 4.$$

$(1 + X_1)^2 + (2 - X_2)^2 > 4$ are the points outside this circle, while $(1 + X_1)^2 + (2 - X_2)^2 \leq 4$ are the points inside and on the boundary of the circle.

c) Suppose that a classifier assigns an observation to the blue class if

$$(1 + X_1)^2 + (2 - X_2)^2 > 4,$$

and to the red class otherwise. To what class is the observation $(0, 0)$ classified? $(-1, 1)$? $(2, 2)$? $(3, 8)$?

- $(0, 0)$ is blue
- $(-1, 1)$ is red
- $(2, 2)$ is blue
- $(3, 8)$ is blue


d) Argue that while the decision boundary in (c) is not linear in terms of X_1 and X_2 , it is linear in terms of X_1, X_1^2, X_2 , and X_2^2 .

By expanding the terms in the formula for the circle, it is apparent that it contains terms like X_1, X_1^2, X_2 , and X_2^2 , in addition to a constant term. This means that the circle can be regarded as linear in terms of X_1, X_1^2, X_2 , and X_2^2 , by allowing these transformations of the variables X_i .

Problem 3

This problem involves plotting of decision boundaries for different kernels and it's taken from Lab video.

```
# code taken from video by Trevor Hastie.
set.seed(10111)
x <- matrix(rnorm(40), 20, 2)
y <- rep(c(-1, 1), c(10, 10))
x[y == 1, ] <- x[y == 1, ] + 1
plot(x, col = y + 3, pch = 19, xlab = expression(X[1]), ylab = expression(X[2]))
```




```
dat = data.frame(x, y = as.factor(y))
```

- (a) Plot the linear decision boundary of the `svmfit` model by using the function `make.grid`. Hint: Use the `predict` function for the grid points and then plot the predicted values $\{-1,1\}$ with different colors.

```
library(e1071)
svmfit = svm(y ~ ., data = dat, kernel = "linear", cost = 0.1, scale = F) # E.g. cost 0.1
```

The following function may help you to generate a grid for plotting:

```
make.grid = function(x, n = 75) {
  # takes as input the data matrix x and number of grid points n in
  # each direction the default value will generate a 75x75 grid
  grange = apply(x, 2, range) # range for x1 and x2
  x1 = seq(from = grange[1, 1], to = grange[2, 1], length.out = n) # sequence from the lowest to the
  x2 = seq(from = grange[1, 2], to = grange[2, 2], length.out = n) # sequence from the lowest to the
  expand.grid(X1 = x1, X2 = x2) #create a uniform grid according to x1 and x2 values
}
x <- as.matrix(dat[, c("X1", "X2")]) # Retrieve only the predictors X1 and X2.
xgrid <- make.grid(x) # Make the grid with the supplied function.
preds <- predict(svmfit, xgrid) # Make predictions on grid of testpoints.
plot(xgrid, preds, col = c("red", "blue")[as.numeric(preds)], cex = 0.2,
     pch = 20)
```



(b) On the same plot add the training points and indicate the support vectors.

(c) The solutions to the SVM optimization problem is given by

$$\hat{\beta} = \sum_{i \in S} \hat{\alpha}_i y_i x_i ,$$

where S is the set of the support vectors. From the `svm()` function we cannot extract $\hat{\beta}$, but instead we have access to $\text{coef}_i = \hat{\alpha}_i y_i$, and $\hat{\beta}_0$ is given as `rho`. For more details see here.


Calculate the coefficients $\hat{\beta}_0, \hat{\beta}_1$ and $\hat{\beta}_2$. Then add the decision boundary and the margins using the function `abline()` on the plot from (b).

Problem 4

Now we fit an svm model with radial kernel to the following data taken from @ESL. Use cross-validation to find the best set of tuning parameters (cost C and γ). Using the same idea as in Problem 4a) plot the non-linear decision boundary, and add the training points. Furthermore if you want to create the decision boundary curve you can use the argument `decision.values=TRUE` in the function `predict`, and then you can plot it by using the `contour()` function.

R-hints:

```
load(url("https://web.stanford.edu/~hastie/ElemStatLearn/datasets/ESL.mixture.rda"))
# names(ESL.mixture)
rm(x, y)
attach(ESL.mixture)
plot(x, col = y + 1, pch = 19, xlab = expression(X[1]), ylab = expression(X[2]))
```



```
dat = data.frame(y = factor(y), x)
```

To run cross-validation over a grid for (C, γ) , you can use a two-dimensional list of values in the `ranges` argument:

```
r.cv <- tune(svm, factor(y) ~ ., data = dat, kernel = "...", ranges = list(cost = c(...),
  gamma = c(...)))
```

For the plot:

```
xgrid = make.grid(x)
ygrid = predict(..., xgrid)
plot(xgrid, col = as.numeric(ygrid), pch = 20, cex = 0.2)
points(x, col = y + 1, pch = 19)
# decision boundary
func = predict(..., xgrid, decision.values = TRUE)
func = attributes(func)$decision
contour(unique(xgrid[, 1]), unique(xgrid[, 2]), matrix(func, 75, 75),
  level = 0, add = TRUE) #sum boundary
```

Problem 5 - optional (Book Ex. 7)

This problem involves the OJ data set which is part of the ISLR package.

- Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.
- Fit a support vector classifier to the training data using `cost=0.01`, with `Purchase` as the response and the other variables as predictors. Use the `summary()` function to produce summary statistics, and describe the results obtained.
- What are the training and test error rates?
- Use the `tune()` function to select an optimal cost. Consider values in the range 0.01 to 10.
- Compute the training and test error rates using this new value for cost.
- Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for `gamma`.
- Repeat parts (b) through (e) using a support vector machine with a polynomial kernel. Set `degree=2`.
- Overall, which approach seems to give the best results on this data?